



# SPEZIELLE THEMEN DER ALGORITHMIK

Patrizio Angelini

Universität Tübingen

[angelini@informatik.uni-tuebingen.de](mailto:angelini@informatik.uni-tuebingen.de)

# INFORMATION

- Lectures: Friday, 10:00 – 13:00 – Room C215
- Assignments: 2
  - At least 1 to participate to the exam
  - Possible bonus points for the exam
- Exam: to be decided


Please register asap to the Moodle (password = **stda**)

# TOPICS

- **Part 1:** General algorithmic techniques for settings in which exact and efficient algorithms are not possible
  - Online algorithms
  - Approximation algorithms
  - Randomization
  - Game Theory
- **Part 2:** More specific algorithmic topics, mainly about graphs
  - Data structures for planar graph embeddings
  - Schnyder realizers for straight-line drawings
  - Euler's formula and Crossing Lemma

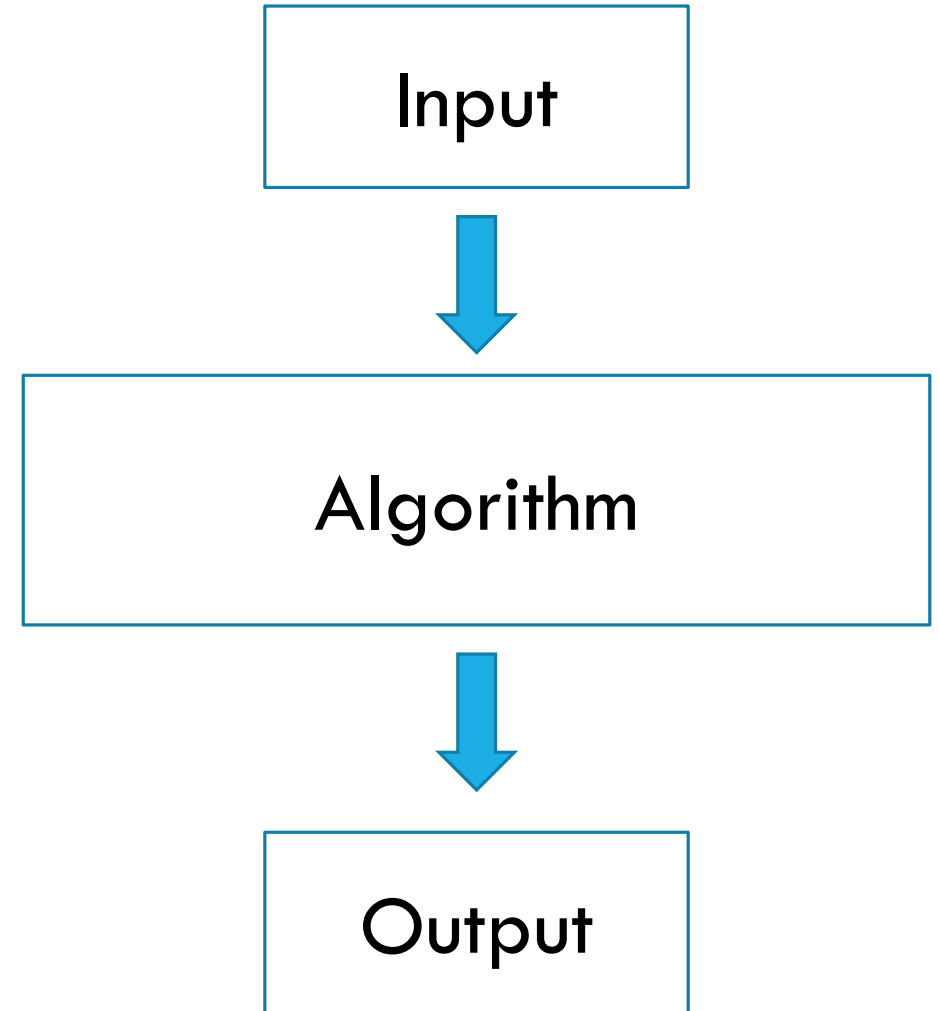
# PART 1

- Given a problem, defining an input and a desired output, we usually aim at designing **exact** and **efficient** algorithms
- **Exact**: correct/optimal solution
- **Efficient**: polynomial running time



# PART 1

- Unfortunately, this is not always possible
- The problem may be difficult
- The input may be difficult to read
- The centralized execution of the algorithm may be not possible



# ONLINE ALGORITHMS

- The input is provided one portion at a time
  - The algorithm has to make choices without knowing the following portions
- The quality is measured by comparing it with an optimum offline algorithm for the same problem and the same input
  - The offline algorithm knows the whole input in advance
- Not to be confused with real-time and streaming
  - Time plays no role here
  - Input has finite size

# COMPETITIVE ANALYSIS

- The measure is performed by considering an **adversary** providing at every step the next portion of the input, which is the worst possible for the current state
  - Corresponds to the worst-case analysis in the offline setting
- **Competitive ratio**: An algorithm is  **$c$ -competitive** if it is at most  $c$  times worse than the optimum

# COMPETITIVE ANALYSIS

- Let  $A$  be an online algorithm for a problem
- Let  $OPT$  be the optimal offline algorithm for the same problem

Algorithm  $A$  is  **$c$ -competitive** if there exists a constant  $\beta$  such that, for every input sequence  $\sigma$

$$A(\sigma) \leq c \cdot OPT(\sigma) + \beta$$

$c$  is called **competitive-ratio**



# APPLICATIONS

- Scheduling
  - Industrial tasks
  - Computer processes
- Load balancing
- Paging
- Routing
- Data structures
- ...
- Ski rental 😊

# SKI RENTAL PROBLEM

- A person  $P$  goes to the Alps and wants to ski
- $P$  does not have any ski, and should decide whether buying or renting them
- Renting costs 1 EUR / day
- Buying costs  $B$  EUR
- The weather is unstable, so  $P$  does not know how many days he can stay
- What should he do?

# SKI RENTAL PROBLEM: OFFLINE OPTIMUM

- The input consists of the sequence  $\sigma$  of good-weather days
- The offline algorithm knows how long is this sequence
  - It knows how many days  $P$  can ski
  - It can decide whether  $P$  should rent or buy
  - *If  $|\sigma| < B$ , then RENT, otherwise BUY*
  - $OPT(\sigma) = \min(|\sigma|, B)$

# SKI RENTAL PROBLEM: ONLINE

- What is the worst-case scenario?
  - How does a bad adversary behave?
- If  $P$  buys ski at day  $d$ , then the adversary makes it rain at day  $d+1$
- If  $P$  does not buy at day  $d$ , then the adversary lets day  $d+1$  be sunny

# SKI RENTAL PROBLEM: ONLINE

- Algorithm 1: *BUY* at day  $d=1$ 
  - What is the competitive ratio?
  - $B/1 = B$
- Algorithm 2: always *RENT* (say, *BUY* at day  $d=\infty$ )
  - What is the competitive ratio?
  - $\infty/B = \infty$
- Algorithm 3: *BUY* at day  $d = B$ 
  - What is the competitive ratio?
  - $(B-1+B)/B = 2 - 1/B$

# SKI RENTAL PROBLEM: ONLINE

- **Theorem:**  $2 - 1/B$  is the best competitive ratio
- **Proof:** Suppose that  $P$  buys at day  $d=d^*$
- The opponent will make it rain on day  $d^*+1$
- If  $d^* < B$ , then c.r. =  $(d^*+B-1)/\min\{d^*, B\} = (d^*+B-1)/d^* \geq 2$
- If  $d^* > B$ , then c.r. =  $(d^*+B-1)/\min\{d^*, B\} = (d^*+B-1)/B \geq 2$

# ANOTHER PROBLEM: K-PAGING

- A two-level memory system, accessed by a CPU, consisting of a **small fast** memory and a **large slow** memory
- Memory is partitioned into **pages** of equal size
- Input: a **sequence** of requests, each specifying a page in the memory
- A request can be served immediately if the page is in the fast memory
- Otherwise, a **page fault** occurs

# K-PAGING

- In case of a page fault, the missing page is loaded from the slow into the fast memory so that the request can be served
- To make room, a page has to be removed from the fast memory
- Which one?
  - A decision must be made without knowing future requests
- The goal is to minimize the number of page faults



# K-PAGING: MOST USED ALGORITHMS

- LRU: Least Recently Used
  - The page to be removed is the one that has been requested least recently
- FIFO: First In First Out
  - The page to be removed is the one that has been added to the fast memory least recently
- How competitive are these two algorithms?

# K-PAGING: MODEL

- Fast memory has  $k$  pages
- Slow memory has  $m \gg k$  pages
- Input: sequence  $\sigma$  of requests
- Cost  $A(\sigma)$  of an algorithm  $A$ : number of page faults

# LRU : LEAST RECENTLY USED

- **Theorem:** LRU is  $k$ -competitive
- **Proof idea:**
  - Divide the sequence  $\sigma$  into **phases** such that in each phase  $\varphi$ 
    - there are at most  $k$  page faults in LRU, that is,  $LRU(\varphi) \leq k$
    - in any (offline) algorithm, even the optimum one OPT, there is at least 1 page fault, that is  $OPT(\varphi) \geq 1$
- This implies that  $LRU(\sigma)/OPT(\sigma) \leq k$ , for any  $\sigma$

# LRU : LEAST RECENTLY USED

- Divide  $\sigma$  into **phases**  $\sigma = \varphi_0 \varphi_1 \dots \varphi_i \dots \varphi_m$  in such a way that:
  - $\varphi_0$  has **at most**  $k$  faults (and at least one)
  - $\varphi_i$  has **exactly**  $k$  faults, for each  $i=1, \dots, m$
- This can be done by following the sequence  $\sigma$  in reverse order, starting from the end, and counting  $k$  faults each time

# LRU : LEAST RECENTLY USED

- Assume that, before applying the two algorithms *LRU* and *OPT*, the fast memory contains the same pages (or is empty)
- In  $\varphi_0$ , the first page fault for *LRU* is also a page fault for *OPT*
  - Hence, *OPT* has at least 1 page fault

# LRU : LEAST RECENTLY USED

- Consider any  $\varphi_i$ , with  $i > 0$
- Let  $p$  be the last page that has been requested in phase  $\varphi_{i-1}$ 
  - Page  $p$  is in the fast memory at the beginning of  $\varphi_i$
- We prove that the set of pages that are requested in phase  $\varphi_i$  contains (at least)  $k$  pages different from  $p$ 
  - This implies there is at least a page fault

# LRU : LEAST RECENTLY USED

- We prove that the set of pages that are requested in phase  $\varphi_i$  contains (at least)  $k$  pages different from  $p$
- Let  $r_1, \dots, r_k$  be the  $k$  requests on which *LRU* had a page fault in  $\varphi_i$
- Let  $r_0$  be the last request in  $\varphi_{i-1}$ , which requested  $p$
- Two cases, based on whether there exist two requests in  $r_0, \dots, r_k$  of the same page, or not
- If not,  $p(r_1), \dots, p(r_k)$  are the pages we are looking for

# LRU : LEAST RECENTLY USED

- We prove that the set of pages that are requested in phase  $\varphi_i$  contains (at least)  $k$  pages different from  $p$
- If yes, let  $r_h = r_j$  be two requests of the same page such that  $r_h, \dots, r_{j-1}$  request different pages
- Since  $p(r_h)$  was the least recently requested page after request  $r_h$ , in order to have a new page fault of the same page  $p(r_h) = p(r_j)$  at  $r_j$ , there are at least  $k$  different pages requested in  $r_h, \dots, r_{j-1}$ 
  - If  $h = 0$ , then these  $k$  pages are different from  $p$
  - If  $h > 0$ , then  $k-1$  of them are different from  $p$ , and  $p(r_h)$  is different from them and from  $p$



# LRU : LEAST RECENTLY USED

- Note that the proof is independent of the optimum algorithm  $OPT$ 
  - It only uses upper and lower bounds
- An analogous argument can be used to prove that algorithm FIFO is  $k$ -competitive
- Can we do better?

# LRU : LEAST RECENTLY USED

- **Theorem:** There exists no  $c$ -competitive deterministic online algorithm for the  $k$ -paging problem such that  $c < k$
- **Proof:**
- Consider a set of  $k+1$  pages.
- For any deterministic algorithm  $A$ , there is an input sequence such that  $A$  has a page fault at every step
  - The opponent asks the (unique) page not in fast memory

# LRU : LEAST RECENTLY USED

- **Theorem:** There exists no  $c$ -competitive deterministic online algorithm for the  $k$ -paging problem such that  $c < k$
- **Proof (continued):**
- For each sequence of  $k$  requests, there is one page that is not present
- When  $OPT$  has a page fault, it can remove from the memory a page that will not be requested in the next  $k-1$  steps
  - In every  $k$  consecutive steps,  $OPT$  pays at most once

# LITERATURE

- S. Albers. *Online algorithms: a survey*. Math. Program. 97(1-2): 3-26 (2003)
- A. Borodin, R. El-Yaniv: *Online computation and competitive analysis*. Cambridge University Press, ISBN 978-0-521-56392-5, pp. I-XVIII, 1-414